

KIRAN RANGAIAH SHASHI

AI Architect & Senior GenAI Engineer | Hands-On Delivery Inside UBS's AI-First Enterprise Strategy | UBS Wealth Management

Application-Layer AI Delivery · Azure AI Factory Engineering · Data-Mesh-Aligned Pipelines · AI Risk Governance · 36-Agent Production Platform

Stanmore, London | 07787 030982 | rskiran79@gmail.com | linkedin.com/in/kiran-rangaiah-shashi-21406b21

EXECUTIVE SUMMARY

UBS runs one of the most coherent AI strategies in global finance: an AI-first mandate built on augmenting human judgement rather than automating it away, a disciplined three-layer value chain (Enabling, Intelligence, Application), and an industrialised 'AI factory' operating model on Microsoft Azure. I don't watch that strategy from the outside — I build inside the Application layer, where the CIO's own framework says the largest monetisation potential sits. Over 18+ years in financial services, including 9+ years working with UBS as contractor and permanent staff combined (3 years permanent), I've personally engineered the RAG pipelines, the 36-agent production platform, and the advisor-facing GenAI tools that put that strategy into working software — provisioning Azure OpenAI capacity, designing data pipelines on data-mesh principles so they're discoverable and trustworthy across business domains, and building the AI risk governance evidence (model cards, bias assessment, HITL gates) that lets any of it survive an FCA audit. I know the difference between a strategy deck and a system in production, because I'm the one who ships the second one.

KEY HIGHLIGHTS

36-Agent Production AI Platform — Application-Layer Delivery at UBS Scale	Azure AI Factory Engineering — PTU Capacity Planning, Agent Service, Prompt Flow	WealthTech Awards 2026 Winner · STAAT AI Suite (7 Production Applications)
Data-Mesh-Aligned Pipeline Design — Domain-Owned, Discoverable Data Products	AI Risk Governance — RiskLab AICE, Model Cards, Bias/Explainability, FCA/BCBS 239	Augmentation-First Delivery — Advisor Co-Pilot Tools, HITL by Design

CORE COMPETENCIES & TECHNOLOGIES

Enterprise AI Strategy Fluency: Application-layer AI delivery within a three-layer (Enabling/Intelligence/Application) investment framework · Augmentation-over-automation design philosophy · AI factory operating model · AI risk governance as a first-class pillar

Azure AI Platform Engineering: Azure AI Foundry (Agent Service, Prompt Flow, Model Catalog), Azure OpenAI Service, Provisioned Throughput Unit (PTU) capacity planning, Azure AI Search (vector + hybrid retrieval), Azure AI Content Safety

Data Architecture: Data-mesh-aligned pipeline design (domain-owned, discoverable, trustworthy data products) · ADLS/Databricks/Unity Catalog · Kafka streaming · PySpark ETL/ELT · Microsoft Purview lineage & classification

Agentic AI & GenAI Engineering: LangGraph, AutoGen, Semantic Kernel, MCP Protocol/API2MCP — orchestrator-specialist multi-agent architecture at 36-agent production scale · RAG (chunking, embeddings, hybrid retrieval) · Python (FastAPI, Flask, full-stack)

AI Risk Governance: Model cards, bias assessment, explainability, AI inventory, RiskLab AICE drift/risk monitoring, policy-as-code, human-in-the-loop approval gates, FCA/BCBS 239 alignment

MLOps & Production Engineering: MLflow, CI/CD (GitHub Actions/Azure DevOps), Terraform/Bicep IaC, Docker/Kubernetes/AKS, APIM (LLM gateway), AI Observability (LangSmith, Braintrust, Ragas)

PROFESSIONAL EXPERIENCE

Associate Director — AI Architect & Senior GenAI Engineer | UBS Wealth Management Apr 2023 – Present · London, UK
Deliver hands-on inside UBS's Application layer — the layer the bank's own CIO framework identifies as carrying the largest monetisation potential in its AI investment thesis — building the augmentation-first tools that put strategy into production.

► Application-Layer AI Delivery — Augmentation, Not Automation

- Built RAG systems end-to-end — ingestion, semantic chunking, embeddings, vector indexing on Azure AI Search, retrieval-augmented pipelines — powering Talk2Docs and enterprise document intelligence, digitising and indexing tens of thousands of internal documents into instantly queryable advisor knowledge
- Engineered LLM application architecture on Azure OpenAI Service, including Provisioned Throughput Unit (PTU) capacity planning to keep pace with UBS's rapidly scaling AI compute footprint — token management, context optimisation, streaming, and Azure AI Content Safety guardrails for regulated deployment

- Designed every AI tool around the augmentation principle at the heart of UBS's public AI philosophy: building co-pilots that make advisors faster and sharper, never systems intended to disintermediate the advisor-client relationship

► 36-Agent Agentic AI Platform — Client Lifecycle AI (OnePass)

- Architected and personally engineered a 36-agent Orchestrator-Specialist architecture on UBS's Nucleus AI Platform, automating AML screening, KYC verification, and client onboarding/offboarding — production-scale proof of the 'AI factory' model's promise of quicker, more agile solution delivery
- Leveraged MCP Protocol, API2MCP Adapters, and the Nucleus Agents Library to deliver multi-step automation with human-in-the-loop gates at every regulatory decision point — auditability by design, not retrofitted
- Consolidated 14+ siloed legacy systems into a single AI-governed client lifecycle framework, reducing onboarding cycle time by ~40% — the kind of operating-leverage gain that feeds directly into the cost/income-ratio improvement the bank's AI monetisation thesis is built on
- Defined reusable AI-platform patterns (RAG pipelines, prompt management, LLM gateway) adopted across 6+ downstream business use cases — the industrialised, repeatable delivery pattern the AI factory model is designed to produce

► STAAT AI Suite — WealthTech Awards 2026 Winner

- Architected seven production applications (STAAT Assist, STAAT Insights, Ask STAAT: CIO, Ask STAAT: BGaP, Source of Wealth, STAAT Schedule, STAAT Email) on a shared infrastructure stack (Airflow, MLflow, LangGraph, Azure OpenAI/Llama, Superset) — advisor-facing tools built in the same augmentation spirit as the bank's flagship advisor co-pilot programmes
- Built proactive-insight and meeting-preparation capabilities that return hours of advisor time per client interaction — the same 'radical efficiency' outcome UBS's own AI strategy identifies as the clearest near-term monetisation path
- Personally engineered the underlying data and ML pipelines: feature engineering, model training/evaluation (MLflow), LangGraph agent workflow design, and Airflow DAG authoring — hands-on from data to deployment across all seven applications
- Delivered measurable advisor productivity gains recognised externally by the WealthTech Awards 2026 — direct evidence that application-layer execution, not just strategic intent, is what separates working AI from a strategy memo

► AI Risk Governance & Regulatory Compliance

- Implemented Responsible AI frameworks — bias assessments, explainability layers, HITL gates, model cards, AI inventory — meeting FCA, BCBS 239, and UBS AI Vision 2030 expectations, giving UBS's public commitment to 'AI risk governance' concrete engineering substance
- Leveraged AI Observability (LangSmith, Braintrust, Ragas) and the RiskLab AICE Framework for drift detection, model performance monitoring (KMPI), and real-time risk alerting — the safeguards a bank operating at UBS's AI ambition needs built in, not bolted on
- Architected an AI model validation and control framework (AI Mart model registry, Compliance & Risk Controls), streamlining self-service deployment on AKS and reducing model-to-production lead times

► Platform Engineering, MLOps & Cloud Infrastructure

- Established MLOps best practice across UBS AI programmes: CI/CD for model deployment, automated RAG evaluation gates, experiment tracking (MLflow), model versioning aligned to production SLAs
- Architected Azure API Management (APIM) as the secure LLM gateway — rate limiting, token quota management, policy-based routing, Private Link/Entra ID zero-trust — the control plane behind every AI service in scope
- Authored Terraform and Bicep IaC for the full Azure AI stack (AKS, Container Apps, Azure AI Search, Cosmos DB, APIM, Key Vault), enabling one-click, governed AI environment provisioning at the scale UBS's Azure-first infrastructure strategy demands

► Team Leadership, AI Strategy & Stakeholder Engagement

- Held six concurrent AI Champion designations across UBS's most strategic AI programmes (UBS GPT, AI Foundations, STAAT, AI Mart, AI Vision 2030, EAIS)
- Worked directly with senior bankers, client advisors, and Wealth Management leadership to translate the firm's AI-first mandate into a prioritised, delivered roadmap — from concept through regulated production
- Mentored data scientists and ML engineers on deployment quality gates and engineering best practice as the AI function scaled

Onshore Performance & DevOps SRE Engineer — AI-Ready Data Foundations (Consulting for UBS) | EPAM Systems

May 2021 – Apr 2023 · London, UK

Built the data and reliability foundation that UBS's Application-layer AI strategy would later be engineered on — before the AI-first mandate had a name internally, the data discipline it depends on was already being built.

- Designed the Genesys Analytics Platform (Greenplum, Apache Airflow, PySpark) as the data intelligence backbone for UBS Wealth Management UK — the precursor architecture to the domain-owned, discoverable data products UBS's later data-mesh strategy formalised firm-wide
- Built ETL/ELT pipelines with full data quality and lineage management, treating each dataset as a trustworthy, reusable product for downstream consumption — the same data-as-a-product discipline central to UBS's Enabling-layer AI investment thesis
- Delivered backend microservices in Python (FastAPI/Flask) exposing PySpark-processed datasets via RESTful APIs, feeding advisor-facing and reporting applications — the plumbing beneath every advisor co-pilot tool that followed

- Integrated Microsoft Purview for automated data classification, lineage, and sensitivity labelling — building the audit-ready evidence trail that AI risk governance at application-layer scale later required
- Architected a Cloud Native Reliability Framework on Azure AKS — automated failover testing, chaos engineering, sub-5-minute MTTR — establishing the production-reliability bar UBS's AI factory model expects of every deployed service
- Prototyped the Greenplum-to-Azure-Synapse/Fabric migration blueprint, positioning UBS's data estate to run on the Azure platform ahead of the bank's stated goal of hosting the majority of its application estate there

Project Manager | Performance DevOps Lead | Data Architect (Consulting for UBS) | Cognizant Worldwide Limited

Feb 2016 – Apr 2021 • London, UK

- Designed Azure Synapse Analytics data warehouse solutions using Data Vault 2.0 — historised, auditable architecture, the evidentiary discipline AI risk governance frameworks later formalised
- Architected Azure Data Factory ETL/ELT pipelines across retail and insurance portfolios with dependency management and self-healing logic
- Led cross-functional engineering teams across UK, Ireland, and India as principal technical consultant across 5+ simultaneous client engagements
- Orchestrated legacy VM monolith migration to Kubernetes (Helm Charts), delivering zero-downtime cutover across mission-critical banking platforms

EARLIER CAREER

Ondot Systems (Fiserv) — Principal Consultant, Performance COE Lead (Dec 2014 – Jan 2016, Bangalore) — core banking load simulation, database optimisation, early Apache Spark adoption.

Sapient • Fidelity Information Services • Mphasis (HP) — Senior Software/Performance Engineer (2007 – 2014) — Java performance engineering, JVM tuning, high-volume systems for e-commerce, payments, and supply chain.

KEY ACHIEVEMENTS & RECOGNITION

- WealthTech Awards 2026 Winner — STAAT AI Suite, seven production applications delivered inside UBS's Application-layer AI strategy
- Built and operate UBS's 36-agent Agentic AI platform — one of the most sophisticated production multi-agent deployments in the firm globally
- Named AI Champion across six concurrent strategic UBS AI programmes — the organisation's highest AI advocacy designation
- Reduced client onboarding cycle time by ~40% and consolidated 14+ legacy systems through governed AI platform engineering
- Top Contributor — UBS GitHub Copilot Prompt Library (2024–2025)

EDUCATION & CERTIFICATIONS

- Bachelor of Engineering (Electronics) — VTU, India
- AI Development with DeepSeek for Developers — Coursera, May 2026 (95%)
- 285+ AI Specialisations — Coursera (ML, Deep Learning, NLP, Cloud AI, Data Science)
- CCNP (Cisco Certified Network Professional) • LoadRunner Expert Certification
- In Progress: Azure Data Engineer Associate (DP-203) • Fabric Analytics Engineer (DP-600) • Azure Solutions Architect Expert (AZ-305) • AWS AI/ML Certification

Right to Work: UK • References available upon request